

Essentials of Computer Science Principles with JavaScript

Kenneth Carter Dodd

2015

Contents

| | | |
|----------|--|-----------|
| 1 | Information | 1 |
| 1.1 | 20 Questions | 1 |
| 1.2 | Communication | 3 |
| 1.3 | Digital Signals | 6 |
| 1.4 | Binary Numbers | 8 |
| 1.5 | Digital Text | 11 |
| 1.6 | Digital Audio and Visual | 11 |
| 1.7 | Error Correction | 12 |
| 1.8 | Encryption | 12 |
| 1.9 | Exercises | 12 |
| 2 | Hardware and Software | 14 |
| 3 | Logic and Binary Arithmetic | 15 |
| 4 | Computation | 16 |
| 4.1 | Sequence of Instructions | 16 |
| 4.2 | Instruction Set | 18 |
| 4.3 | Abstraction | 19 |
| 4.4 | Conditional Instructions | 20 |
| 4.5 | Flowchart | 21 |
| 4.6 | Iteration | 23 |
| 4.7 | Sequence, Condition, Iteration | 24 |
| 5 | Source Code | 25 |
| 5.1 | Programming Languages | 25 |
| 5.2 | JavaScript | 25 |
| 5.3 | HTML and CSS | 26 |
| 5.4 | Style and Conventions | 26 |
| 5.5 | Organization and Version Control | 26 |

| | | |
|-----------|---|-----------|
| 6 | Variables | 27 |
| 6.1 | Valid Variable Names | 27 |
| 6.2 | Reserved Words | 27 |
| 6.3 | Primitive Data Types | 27 |
| 7 | Operators | 28 |
| 7.1 | String Concatenation | 28 |
| 7.2 | Arithmetic | 28 |
| 7.3 | Boolean Algebra | 28 |
| 7.4 | Comparisons | 28 |
| 8 | Control Structures | 29 |
| 8.1 | if-else | 29 |
| 8.2 | nested if-else | 31 |
| 8.3 | while | 35 |
| 8.4 | for | 37 |
| 9 | Functions | 38 |
| 9.1 | Inputs and Outputs | 38 |
| 9.2 | Function Scope | 40 |
| 9.3 | Recursion | 42 |
| 9.4 | Exception Handling | 43 |
| 9.5 | Side-effects | 45 |
| 9.6 | Closure | 45 |
| 10 | Data Structures | 48 |
| 10.1 | Objects | 48 |
| 10.1.1 | Code Exercises | 50 |
| 10.2 | Arrays | 50 |
| 10.3 | JSON | 52 |
| 11 | Constructors | 53 |
| 11.1 | new Constructor | 53 |
| 11.2 | Inside a Functional Constructor | 53 |
| 11.3 | The specs Object | 54 |
| 11.4 | Extending Objects | 58 |
| 11.5 | Copy Constructor | 60 |
| 11.6 | Getters and Setters | 60 |
| 11.7 | | 60 |
| 11.8 | Sharing Hidden Information | 60 |

CONTENTS

iii

12 Searching and Sorting

63

Chapter 1

Information

We use the word information frequently. Think of an example that you think contains information. Perhaps you thought of a book or a website. Can you identify what exactly about these things contain information, or how you might quantify an amount of information it contains?

1.1 20 Questions

If you are not familiar with it, 20 Questions is a game where one person thinks of an object which other people must guess. The person who knows the object may be asked any yes or no question until 20 such questions have been answered. The premise of the game is that if the questions are clever enough, only 20 questions are needed to guess what it is. But, if the questions are not good then it may not be enough to guess the object. But think of how many objects actually exist; could 20 questions really be enough to narrow it down?

Lets start with a simple example and limit the possibilities to the letters of the English alphabet. There are 26 letters. How many yes or no questions are required to guess one of the letters? One could adopt the strategy of asking questions like "is it the letter A?". This is a gamble. On the one hand, the letter could be guessed on the first try if the answer is 'yes'. But this type of question only eliminates one possibility at a time if the answer is 'no', and so it could take up to 25 questions to eliminate all but one of the letters.

```
                ABCDEFGHIJKLMNOPQRSTUVWXYZ
Guess A:  BCDEFGHIJKLMNOPQRSTUVWXYZ
Guess I:  BCDEFGH JKLMNOPQRSTUVWXYZ
```

```

Guess E:  BCD FGH JKLMNOPQRSTUVWXYZ
Guess Y:  BCD FGH JKLMNOPQRSTUVWXYZ
Guess X:  BCD FGH JKLMNOPQRSTUVW  Z
Guess G:  BCD F H JKLMNOPQRSTUVW  Z
Guess Z:  BCD F H JKLMNOPQRSTUVW
Guess Q:  BCD F H JKLMNOP  RSTUVW
Guess P:  BCD F H JKLMNO   RSTUVW
Guess V:  BCD F H JKLMNO   RSTU  W
Guess U:  BCD F H JKLMNO   RST   W
Guess W:  BCD F H JKLMNO   RST
Guess O:  BCD F H JKLMN    RST
Guess M:  BCD F H JKL  N    RST
Guess S:  BCD F H JKL  N    R  T
Guess T:  BCD F H JKL  N    R
Guess R:  BCD F H JKL  N
Guess B:   CD  F H JKL  N
Guess D:   C   F H JKL  N
Guess K:   C   F H J  L  N
Guess C:         F H J  L  N
Guess J:         F H    L  N
Guess L:         F H      N
Guess N:         F H
Guess F:         H

```

We could ask questions like "does the letter come before the letter N?". If the answer is 'yes', then half of the letters are eliminated and half remain as possibilities. If the answer is 'no', then the other half are eliminated and half remain. The 13 remaining letters can be divided again, and again eliminating half of the letters with each 'yes' or 'no' answer until only one letter remains. So, at most 5 answers are needed to guess the letter.


```

                ABCDEFGHIJKLMNOPQRSTUVWXYZ
Guess < N: ABCDEFGHIJKLM
Guess < G:           GHIJKLM
Guess < J:           GHI
Guess < H:           HI
Guess < I:           H

```

These relationships can be represented by what is called a binary tree. The 'root' of the tree is at the top. From each node in the tree there are

two possible ways to go further down the tree. Each time the tree splits this represents a yes/no or left/right, or whatever the *binary* difference may be which is used to exclude the other branches of the tree. Looking at this figure it can be seen that any letter can be guessed in either 4 or 5 guesses.



This kind of strategy seems to rely on a way to compare the letters to see if they come 'before' or 'after' another letter in order to properly answer the question "does the letter come before the letter N?". But we could have instead asked "is the letter among one of the letters A, B, C, D, E, F, G, H, I, J, K, L, or M?" This may seem like cheating, but this still qualifies as a single yes or no question with a single response.

What we are doing is eliminating groups of letters at a time, and we are free to eliminate any number at a time. That is why this game works for a much broader set of objects than just those that can be ordered like letters or numbers. In the game, questions can be things like "is it a person?" or "is it made of metal?".

The real reason these kind of questions work is that every object has some defining properties that can be used to differentiate it from all other objects, and the premise of the game is that any property can be reduced to a 'yes' or 'no' question. We can think of each object as either having a property, or not having it, which gives the answer to any possible question. How we differentiate things from each other is the essence of information. If we can't tell any difference, then there is no information.

1.2 Communication

Information is defined by being able to rule out alternatives. A statement contains no information in itself unless we also know how the statement could

have been different. Saying that it is raining outside only contains information because we have alternatives to it raining, such as a clear sky or snowing. The alternatives don't have to necessarily be logical or possible in reality.


For instance, we can say that the formula of water is H_2O . This contains information even though we do not think that water actually could have another formula. This is to differentiate the formula of water from some other substance, for example ethanol which is C_2H_6O . It also does not rule out some other substance having the same formula. For instance diethyl ether is also C_2H_6O , but it has a different structure which gives it different chemical properties. More information is needed to differentiate chemicals with the same formula.

To explore the idea of information we will look at forms of communication: how to send information from one place to another. There are several forms of communication based only on the ability to send what is called a binary signal. The word binary basically means a set of two, or a pair, like in the word bicycle for two wheels. Here a binary signal means a signal that only has two possible states. When we receive one state, the only information we have gained is that we didn't get the other state.


Examples of this might be a smoke signal, or a light, or an electrical signal that either on or off. All we can tell is if we can see the signal, or we can't. This is exactly like a 'yes' or 'no'. So, we know we could receive answers to play the 20 questions game by this method, but how could we ask the questions if all we can do is say 'yes' or 'no'?

In this case, we seem to need some pre-arrangements. Suppose ahead of time we worked out a question that had 'yes' or 'no' answer, so that when we are sending the signal they know what question we're answering. This doesn't seem very useful because we would need a way to change which question we're answering for each signal.

Since we have the ability to turn the signal on and off over time, we could make longer or shorter signals to mean different things. A short signal could be a 'yes' to one question, and a long signal could be a 'yes' to a different question. Similarly we could think of short or long pauses without a signal as being a 'no' to different questions.



In Morse code, a sequence of short and long signals are used to specify a letter of the alphabet. Short signals (dots) and long signals (dashes) tells us which way to go in the tree of possible letters, and rules out the alternatives, with short pauses in between each signal. A long pause is used to indicate that a new letter is being started.



For example, suppose the first thing we receive is a short signal. In the above tree, you would start at the very top circle, called a node of the tree. If the first signal is a dot (short), then the letter is somewhere in the left side of the tree (E, I, A, S, U, R, W, H, V, F, L, P, or J). If the first signal is a dash, then it must be in the right side of the tree.

The next thing that will happen is to have either a short pause or a long pause. If it is a short pause, then we know we need more signals to know what letter it is. If it is a long pause then we have all the signals and we know which letter it is. So, if we receive a short signal and then a long pause, then that means the first letter is an 'E': •.

The letter 'X' is in the right hand side of the tree, so the first signal will be a dash. From there it is two dots and then another dash to get to the 'X' branch of the tree: —••—. Notice that any letter can be specified by as little as 1 signal, and no more than 4 signals. It seems like this requires

□ ● □ ● ● □ □

message: ●● ■● ■ ● ●■ ■●● ●●●● ●● ■● ■■●
■■■ ■ ●●●● ● ●■ ●●● ■■■ ● ■ ● ●■ ■●■● ●●●●
■■■ ●●■ ●■● ●●●● ● ●■●● ●●●■ ● ●●●

1.3 Digital Signals

Morse code is an example of a digital signal, but the idea can be generalized to send much more information than letters of the alphabet. Morse code was created at a time when the encoding, transmission, and decoding had to be done manually, and was limited by human speed. Modern technology is able to do these functions much more rapidly, and uses digital signals to send text, audio, video, and other forms of information for cellphones, television, and the internet. So, how does all that information get broken down, if not with dots and dashes?

The way this is done is very similar to how sheet music is written. In music, individual notes are supposed to be played for a specific amount of time, which is regulated by a beat frequency. Music can be played faster or slower by making the beats shorter or longer, but the relation between how long each note lasts is fixed by the sheet music.

In a digital signal, if we know the beat frequency all we need to do is look for *changes* in the signal. Even if the signal does not change we can gain information because it *could* have changed in the same way that every beat *could* be a different note.

We already used this kind of scheme with the Morse code by having long versus short signals, and we can be specific about how many beats each element lasts:

| Morse code | # of beats |
|------------------------|------------|
| dot | 1 |
| dash | 3 |
| pause between elements | 1 |
| pause between letters | 3 |
| pause between words | 7 |

Using these rules we could break up the letter 'C' into 11 individual beats from the code $\text{---}\bullet\text{---}\bullet$. In the following image each box represents one beat. If the box is filled in that represents a signal, and if it is empty then there is no signal.



If we loosen the restrictions on how long each signal or pause has to last, allowing them to last any number of beats, then the total number of different possible sequences is much larger for a given number of beats. If any beat can have either value, then it is called a bit. Bit is short of binary digit. Since this communication is broken down into individual binary digits we call it digital.

Each bit has two possible values, and so doubles the number of total possibilities. Two bits have four possible ways they could happen. Three bits have 8 ways. Eleven bits has $2^{11} = 2048$ different ways, and so on.

I could set the signals with a regular pattern of bits, perhaps to tell the receiver what the beat frequency is.



Or I could set the bits randomly, such as with the flips of a coin.



A digital signal is an abstract idea, and so can be implemented by different hardware using different mediums, and converted from one to another. When you load a web page, that stream of bits originated as a file on a hard-drive which was probably stored by magnetic fields on a metal disk. Then converted to electrical pulses when loaded into the servers memory where it was stored by transistors and capacitors.

It's then converted again and sent into the network electronically over wires to a service provider. It may then be converted into pulses of light which travel through optical fibers until it reaches another large service provider. Or it is converted to radio waves and transmitted to a satellite which does additional conversion from electrical to radio to send it to another part of the planet. Finally the process is complete bringing it into your computers memory where it is once again stored by little capacitors. Your computer then converts the bits into images and sounds through your browser so that you can experience it on your monitor or speakers.

However, we don't have to worry about *how* the information is physically being stored or transmitted. The low level hardware is abstracted away so that we only have to deal with the idea of bits.

1.4 Binary Numbers

The binary number system can be analyzed in several different ways. In one view, the concept of a particular number exists apart from the way we write it, and decimal versus binary are just two ways of doing it. The number 5 is a concept apart from the symbol '5'. So is the number '6'. But just like with letters, we need a way to tell which number we're talking about. One way to do it is to have a separate symbol for every number, but that is not practical.

There is no single way to represent a number, and in fact Morse code has its own way of representing numbers. But we will use a standard representation that closely resembles the decimal system that is useful for doing arithmetic. Instead of basing a digit on a power of ten ($1 = 10^0, 10 = 10^1, 100 = 10^2, 1,000 = 10^3 \dots$), we base it on a power of 2 ($1 = 2^0, 2 = 2^1, 4 = 2^2, 8 = 2^3, 16 = 2^4 \dots$). In binary there are only two symbols: zero (0) and one (1), and represent the binary digits (bits). This is also called base-2, as opposed to base-10, since the 'base' of the exponent used is 2.

| power | decimal | binary |
|-------|---------|---------|
| 2^0 | 1 | 1 |
| 2^1 | 2 | 10 |
| 2^2 | 4 | 100 |
| 2^3 | 8 | 1000 |
| 2^4 | 16 | 10000 |
| 2^5 | 32 | 100000 |
| 2^6 | 64 | 1000000 |


Like decimal numbers, the least significant digit (the 1's place) starts on the right, and more significant digits go toward the left. To write the number 2, a 1 is put in the 2's place, and a 0 in the 1's place: $2_{(10)} = 10_{(2)}$. I am also using the subscript (10) and (2) to separate base 10 from base 2 numbers. The number 2 looks like 10 in binary because the 1 is in the 2's place, and every zero to the right represents a power of 2, like the zeros represent powers of 10 in the decimal system.

Converting a number from binary is fairly straight forward as long as you know the numeric value of each place. If there is a 1 in that place then you add it to the total, and if there is a zero then you don't. For example, the number 1011 has a 1 in the 1s, 2s, and 8s place so we add them together to get $1+2+8 = 11$. The number 1111 has all four which gives $1+2+4+8 = 15$. And notice that adding them all together is 1 less than the value of the next bit, which is 16.


Converting a number to binary starts with the most significant bit that has a value less than the number you want to write in binary. To write the number 37, we start by finding the largest power of two that is less than 37. 64 is too big, so we start with 32. That means the 32s bit is set to one. The remainder is 5 ($37 - 32$), and so we have to set the next biggest bit that is smaller than 5. 16 and 8 are too big so the next bit is in the 4s place, which

leaves a remainder of 1 ($5 - 4$). 2 is too big so it has to go in the 1s place (which makes sense). Writing the bits out with zeros for the bits that didn't get used gives: 100101. Try converting it back to decimal to verify it gives the correct number.

Like with Morse code letters, we can construct each number based on a yes/no binary method. If we start with the least significant bit, it has two possible values: 0 or 1. Once that bit is determined by moving down one branch of the tree, the second bit has two possible values. Every additional bit has two possible values, doubling the possible number of numbers.



You may notice that there are not as many numbers in this tree as there are letters in the Morse code tree. That is because binary numbers are not designed to take advantage of a variable length number of bits. Adding another zero most significant bit does not change the number like adding additional dots did for letters. 0100 is the same number as 100. Converting to/from binary can be thought of as traversing the tree to the number in question. Remember that the top of the tree is the *least* significant bit. So, 1101 leads to the number 13.



While individual bits carry information, one can see that quite a few bits are required to represent anything useful for us humans. A byte is a short sequence of bits, and is the smallest 'chunk' of information a computer tends deal with at one time. It is currently widely accepted to mean exactly 8 bits,

although historically the size of a byte has varied.


8 bits have a total number of $2^8 = 256$ possible values. We can take a continuous stream of bits and group every multiple of 8 together into a continuous stream of bytes. This makes using the stream of information a little easier since each 'chunk' can differentiate between 256 possibilities, instead of just 2.


If we are limited to 1 byte, then a binary number has to have exactly 8 digits. The number 0 would be 00000000, the number 1 would be 00000001, and so on. The most significant bit has a value of $2^7 = 128$, and so the byte 10000000 in binary is equivalent to the number 128. If all the bits are set, 11111111 has a value of $2^7 + 2^6 + 2^5 + 2^4 + 2^3 + 2^2 + 2^1 + 2^0 = 255$. So a byte can have integer values from 0 up to 255.

Bigger numbers of course require more bytes to be represented. We will deal with how to represent negative numbers and fractions in a later chapter. However, the basic idea is that we are limited by how many possibilities can be represented.


1.5 Digital Text

1.6 Digital Audio and Visual





1.7 Error Correction



1.8 Encryption

1.9 Exercises

1. A single flashlight has two possible states that can be perceived by a person at a distance (on and off). How many different states could be perceived using two flashlights? Assume the flashlights are identical. If they are held side-by-side, could someone tell which one was on, when only one of them is on? If someone is far enough away, will the two flashlights appear as a single flashlight when they are both on?
2. Suppose a red filter is added to one of the flashlights, and a blue filter is added to the other one. How will having two different colors change how many possible states can be perceived by a person at a distance

away? If the two lights look like a single light from far away, what color will they appear when both flashlights are on?

3. Suppose we have a 4 state transmission system labeled by the following digits: 0 (off), 1 (on), 2 (on), 3 (on). Come up with a code that uses the off state to separate signals, the letters of the alphabet are specified by a sequence of digits, and a space is also a sequence of digits. What is the minimum sequence of digits needed to send any of the 26 letters or a space?

Chapter 2

Hardware and Software

Chapter 3

Logic and Binary Arithmetic

Chapter 4

Computation

What do we mean by computation? In a very general sense, a computation is a goal that is possible to reach by some number of intermediate steps. You may be familiar with doing computations in math, such as multiplying two numbers. But the idea of computing is much broader than your math class. Any problem that *can* be broken down into specific steps to reach a specific conclusion is called computable. The details of the exactly what steps to take in a computation is called an **algorithm**. Figuring out how to combine things that you *can* do to accomplish something you can't do yet is the essence of computer science.


4.1 Sequence of Instructions

For example, think about how you follow directions when traveling from one location to another. The first time you go somewhere you probably need to either ask someone how to get there, or follow the directions of a GPS device or a map. The ultimate problem is to go from A to B. Directions might come in the form of turning left or right on a particular street, then following that street for some distance, then making another turn, and so on.

Solving the problem of going from A to B can be broken down into some number of simpler problems. You don't need anyone to explain how to turn left/right or follow a street because you already know how to accomplish those tasks. All you need to know is what order, or sequence, to do them in to get to where you're going.

A electrical/mechanical computer starts with a very limited number of

instructions that it can understand, and 'knows' how to accomplish. What we have to do is to figure out how to accomplish more complicated actions by breaking them down into a larger number of simpler ones until we are *only* telling the computer those that it can understand.



Come up with a sequence of instructions to give someone that doesn't have the map. They will start at A, and want to arrive at B. As an example, consider these instructions starting at A. Which location would you arrive at?

Go north.

Turn right when you reach Market St.

Go several hundred feet then turn right again.

Go down the street and it's on the left.

Is the set of instructions specific enough to get there? Would the instructions be different if the person didn't know which way was North, or was a poor judge of distance? How can the instructions be altered to be more robust?

4.2 Instruction Set

One issue with the problem of giving someone a sequence of instructions is that we need to know a little more about what they know how to do, so that we know what kind of instruction we need to give them, and how specific we need to be.

Before we start listing the sequence of instructions, we should write down what instructions we have to work with: the instruction set. Suppose we came up with the following set of instructions to choose from, and using the # symbol to mean we can use that instruction with any number substituted in for #.

Directions Instruction Set

- Turn Left.
- Turn Right.
- Go Forward # feet.

What would the directions look like now if we could only use some sequence of these instructions? Lets assume the person will start out facing the road, wherever they begin, and let's tell them how to go from A to B using only these instructions. One sequence might be the following.

Turn Right.

Go Forward 200 feet.

Turn Left.

Go Forward 200 feet.

Turn Right.

Go Forward 100 feet.

We could make this more compact by coming up with codes to represent each instruction, along with a definition of what that code means. Instead of using the # sign to represent any number, I am going to list what is supposed to represent this number with what is called a variable inside parenthesis of the codes. If there is no variable for that instruction, then there is nothing listed inside the parentheses. When the code for instruction is used we can

replace the variable with the number we want to use for it.

Directions Instruction Set Codes

L() : Turn Left.

R() : Turn Right.

U() : Make U-turn.

F(x) : Go Forward x feet.

Now the sequence of instructions would look like the following.

R()

F(200)

L()

F(200)

R()

F(100)

4.3 Abstraction

Think about this. Once they know how to go from A to B, then you don't need to tell you how to get there anymore. Suppose now they need to go to location C. You could tell them to go from A to B first, then give further instructions of how to go from B to C. We can alter the instruction set by defining a new instruction to go from A to B using the codes for the other instructions that they already know.

AB() : go from A to B using the following

R()

F(200)

L()

F(200)

R()

F(100)

Now I can write a sequence of instructions to go from A to C using this new instruction AB(). The details of how to get from A to B are now *abstracted*, which makes writing more complicated directions easier.

AC() : go from A to C

AB()

F(100)

L()

F(500)

4.4 Conditional Instructions

So far we have created instructions that are intended to be followed, in order, no matter what. But in practice we know that not everything can be described quite that simply. What use would the instructions be if the road was blocked off for construction? We may not know that ahead of time, so we need to be able to prepare our instructions to handle that possibility.

The idea of a condition is that one set of instructions will be followed if the condition is true, but a different set of instructions should be followed if it is false. Here the condition might be, "Can you move forward?". If they can, then they can move forward down the road, but if they can't we need to give alternative instructions of what to do from that point on.


We can come up with a codes for following conditional instructions, just like for other instructions. However, instead of the codes telling us where to go in the map, it tells us where to start reading other instruction codes. The IF-ELSE code says that if a condition is true, the instructions immediately following the IF code are followed. But if the condition is false, the instructions immediately following the ELSE code are followed instead.

Suppose we give the same instructions as above, but we know that sometimes Broadway St. is sometimes closed to traffic. We can build in a condition that if when the driver gets to Broadway and it's closed, they can go an alternate rout. The condition is "Is Broadway St. open to traffic?". If the answer to that is 'yes', we follow the normal instructions. But if it is 'no', then we follow the instructions under the ELSE code until the ENDIF code is reached.


```
AC() : go from A to C
    AB()
    F(100)
    IF : Is Broadway St. open to traffic?
        L()
        F(500)
    ELSE :
        U()
        F(200)
        R()
        F(300)
        R()
        F(500)
        L()
        F(200)
    ENDIF :
```

4.5 Flowchart

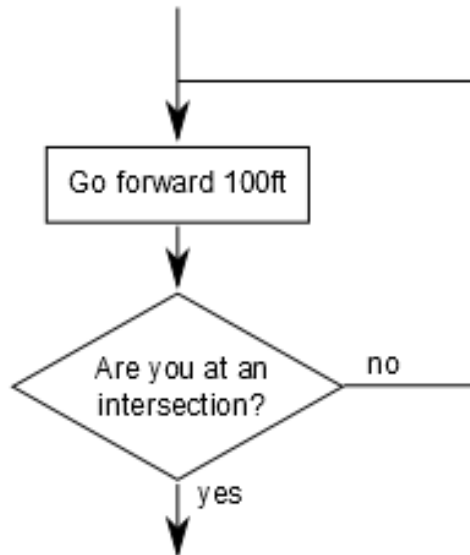
A flowchart is a way of organizing this conditional nature of the directions in a visual way. Here, boxes are the individual directions, and the diamond is the condition which has two possible set of directions. Visualizing the instructions this ways shows that an algorithm itself is like a map, with more than one possible path depending on what is true during any given time through the algorithm. Describing the algorithm using codes or the flowchart are made completely analogous to each other. Flowcharts can be used to plan out an algorithm before implemented with a specific set of codes.




4.6 Iteration

If a conditional statement results in repeating a previous set of instructions, it is called an iteration. This can be visualized using a flowchart. Suppose instead of having to measure the distance needed to drive after each turn, we want to tell the driver to simply drive until an intersection is reached.

We can have a conditional block asking whether or not they are at an intersection. If not, then drive forward some more. If they are, then make the required turn. This will result in what is called a loop. "Are you at an intersection?" serves as the loop condition, which causes the loop to keep repeating until the driver arrives at an intersection.



Perhaps we would like to check to see if they are at an intersection before driving forward, instead of afterward. The main difference is that if the driver is already at an intersection, they can go ahead do the next instruction without having to drive forward (which might cause them to miss the turn).



4.7 Sequence, Condition, Iteration

Any algorithm can be described using some combination of these three basic control structures. When creating a programming language, we must ensure that these functions are possible to describe and implement within the language so that any algorithm can be written using that language. To make the set of codes for driving directions complete, we really need a way to do iteration. The WHILE code will be used to do this.

As long as the condition associated with the WHILE statement is true, the instructions immediately following the while statement are followed, and then the condition is reconsidered. When the condition is false, the instructions under the WHILE statement are no longer followed, and instead the instruction after the ENDWHILE code is followed.

The while loop could be used to define a generic 'forward' code. Since we want it to repeat when the driver is not at the intersection, we have to change the question slightly.

F() : go forward until intersection is reached.

WHILE : Are you not at an intersection?

F(100)

ENDWHILE :

This is how programming languages start to be developed. Once some basic operations, or instructions, are built in to the language, many other and more complicated operations can be built using them.

Chapter 5

Source Code

Source code is text that is readable by a human, but that can be translated into operations that a computer knows how to perform. Source code is always written in a plain text file, and can be edited by any plain text editor.

5.1 Programming Languages

A programming language is one particular way of representing and organizing a program at the human readable level in source code. It is an abstraction of the underlying operations that the computer will perform when the program is executed. In some sense, a programming language is arbitrary since there is usually more than one way to think about or describe the same thing. But since we require a programming language to be able to represent any program, there is a minimum set of concepts it must be capable of expressing.

Before a program can be executed, the source code must be translated into those operations the computer understands. Fortunately you don't have to do that step. The idea of a programming language is that it is specific enough that another program can be written to do the translation for you. That program is either called an interpreter, or a compiler, depending on when and how the translation step is done.

5.2 JavaScript

Although this book focuses on learning the JavaScript language, the larger idea is to understand the underlying concepts because they appear in other languages. Instead of thinking of yourself as becoming a JavaScript programmer, you should try to think of yourself as becoming a general programmer.

Once you learn one language, and understand the underlying ideas, learning another language should be easier because you know what to look for.

Just as different programming languages express the same concepts in different ways, JavaScript itself contains different ways to express the same concepts. This is because of how JavaScript has evolved since its creation. Programs can be written entirely using a subset of the language. You might even consider different sub-sets as dialects of JavaScript because the programs will look different, even though they are all written in the same overall language.

This book is not meant as a comprehensive course on JavaScript and all the different features it contains. We will focus on one particular sub-set of the language and a way of writing programs. If you decide to use code written by other programmers, just remember that they may not write their programs the same way we have.

5.3 HTML and CSS

HTML stands for HyperText Markup Language. There is some debate as to whether HTML qualifies as a programming language. What we mean by programming language has to do with what it is able to compute, or what algorithms it is able to implement. HTML cannot be used to compute things in an obvious way, so for our purposes we may not consider it as a programming language.

However, we do use it to specify information in a web browser in a human readable fashion, such as what files to load and how to display things, which the browser then interprets into the operations it needs to accomplish those tasks. Often accompanying HTML is CSS, which stands for Cascading Style Sheets. CSS is an additional notation used to help specify the display of content in the browser.

5.4 Style and Conventions

5.5 Organization and Version Control

Chapter 6

Variables

6.1 Valid Variable Names

6.2 Reserved Words

6.3 Primitive Data Types

- Boolean
- Null
- Undefined
- Number
- String
- Object

Chapter 7

Operators

7.1 String Concatenation

7.2 Arithmetic

7.3 Boolean Algebra

7.4 Comparisons

Chapter 8

Control Structures

To make our code capable representing any computation, certain control structures are required. We need a way of running a particular command conditionally on some other information. And we need a way of executing the same code more than once.

8.1 if-else

Conditional execution in JavaScript is implemented using the if-else structure, or just called an if-statement. To use an if-statement I first need a quantity that reduces to a boolean value: either true or false. This value becomes the condition under which a particular piece of code is executed if it is true. The condition is placed between the set of open and closed parenthesis. Immediately following the parenthesis is the piece of code that will be executed if the condition is true. The code is grouped together using an open-closed set of curly brackets.

```
1 var a = true;
2 var b = false;
3
4 if (a) {
5     console.log("This will be printed.");
6 }
7
8
9 if (b) {
10    console.log("This will not be printed.");
11 }
```

I can also specify that something is done if the value is *not* true, which will fall under the `else` part. Not true is the same as saying false.

```
1 if (a) {  
2     // nothing done here  
3 }else{  
4     console.log("This will not be printed.");  
5 }  
6  
7  
8 if (b) {  
9     // nothing done here  
10 }else{  
11     console.log("This will be printed.");  
12 }
```

Even though the above example is completely valid, it can be simplified since there are no operations done in the top part of the if-else. I can use the logical 'not', `!`, to negate the boolean values. True becomes false, and false becomes true. This flips the roles of the if-else blocks, and so I can get rid of the 'else' part, while keeping the same logic. That is, this version of the control structure is equivalent to the previous one.

```
1 if (!a) {  
2     console.log("This will not be printed.");  
3 }  
4  
5  
6 if (!b) {  
7     console.log("This will be printed.");  
8 }
```

In the following example, I am testing to see if a particular number is even or odd, and I want it to print something different in each case. To see if a number is even, I used the modulus operator. If the number divided by 2 has a zero remainder, then the modulus will return zero, and I know that 2 is a factor of the number. The `===` operator returns true if `myNumber % 2` is zero, and false if it is not zero. You should always give an if-statement a boolean value, and so comparison operators can be used to convert numbers to booleans. The comparison has the lowest order of operations, and so it will be done last.

```
1 var myNumber = 321;  
2  
3 if (myNumber % 2 === 0)  
4     console.log(myNumber + " is even.");  
5 else  
6     console.log(myNumber + " is odd.");
```

All of the operations together, `myNumber % 2 === 0`, becomes the condition of the if-statement. Only the first thing following the condition is associated with the 'true' case. The piece of code immediately following the else is not executed if the condition is true, but it is if the condition is false. If you wish to have multiple commands associated with the condition, curly braces are used to group the code together.

```
1 var myNumber = 321;
2 var halfInteger;
3
4 if (myNumber % 2 === 0) {
5     console.log(myNumber + " is even.");
6     halfInteger = myNumber/2;
7 } else {
8     console.log(myNumber + " is odd.");
9     halfInteger = (myNumber - 1)/2;
10 }
11
12 console.log(halfInteger + " is half integer.");
```

Here I not only print out whether the number is even, but divide it by two if it is even. If it is odd, I subtract 1 from the number before dividing so that the result is still an integer. If I had not put curly braces around each set of code, then the interpreter would not have understood that I wanted the second statement in each block to only be executed for that condition.

8.2 nested if-else

Any code can be placed in an if-statement, including more if-statements. A common problem is that a variable could have several possible values, and something different should happen for each value.

```
1 var first = "hello";
2 var second;
3
4 if (first === "foo") {
5
6     second = "bar";
7
8 }else if (first === "hello") {
9
10     second = "world";
11
```

```
12 }else{
13     second = "idk";
14 }
15
16 console.log(first + " " + second);
```

The first if-statement checks for one value. If it is that value it does the thing it should for it. If it's not that value, it goes to the **else** part. The only thing in the else of the first if-statement is another if-statement that checks for another value. If it's the second value it does that thing, but if it's not it goes to the else of the second if-statement.

However, any structure of nested if-else statements is valid.

```
1 var clouds = true; // is it cloudy outside
2 var rain = false; // is it raining outside
3 var daylight = true; // is it during the daytime
4 var lights = false; // are the lights on in the house
5
6 var activity;
7
8 if (daylight) {
9
10     if (clouds) {
11
12         if (rain) {
13             activity = "play video games.";
14         }else{
15             activity = "play basketball.";
16         }
17
18     }else{
19         activity = "go swimming.";
20     }
21 }else{
22
23     if (lights) {
24         activity = "work on homework.";
25     }else{
26         activity = "go to sleep.";
27     }
28 }
29
30 console.log("I think I'm going to " + activity);
```

What will be printed above given the values of `daylight`, `clouds`, and `rain`? Does the value of `lights` affect what's printed, or when would it?

What would the values have to be in order for them to decide to "go swimming"?

Certain logic can sometimes be expressed in more than one way. Boolean algebra can be used to rewrite if-else structures into a series of equivalent if-statements using the logical operators.

```
1 var clouds = true; // is it cloudy outside
2 var rain = false; // is it raining outside
3 var daylight = true; // is it during the daytime
4 var lights = false; // are the lights on in the house
5
6 var activity;
7
8 if (daylight && clouds && rain) {
9
10     activity = "play video games.";
11
12 }else if (daylight && clouds && !rain) {
13
14     activity = "play basketball.";
15
16 }else if (daylight && !clouds) {
17
18     activity = "go swimming.";
19
20 }else if (!daylight && lights) {
21
22     activity = "work on homework.";
23
24 }else if (!daylight && !lights) {
25
26     activity = "go to sleep.";
27 }
28
29 console.log("I think I'm going to " + activity);
```

While this is logically equivalent to the first structure, it checks and rechecks the same values at each if-statement. It is also harder to see which scenarios, or cases, have a prescribed activity. And what if a `!` was accidentally omitted? While the choice of which types of structures to use is up to you, your choice should be guided by both the purpose of the structure, and how easy it is to understand and debug if it doesn't work as expected.

Let's look at exactly how an if-else can be re-written with logical operators, or vice-versa. A single 'and' operation can be thought of as a nested

if-else statement where both conditions have to be true for the result to be true. If either condition is false (or both), the result is false. The two if-else structures in the following example are equivalent.

```
1 var a = true;
2 var b = false;
3 var result;
4
5
6 if (a) {
7
8     if (b) {
9         result = true;
10    }else{
11        result = false;
12    }
13 }else{
14     result = false;
15 }
16
17 if (a && b) {
18     result = true;
19 }else{
20     result = false;
21 }
```

For an 'or' operation, if at least one of the conditions is true (or both), then the result is true. This can be represented by a series of if-else statements. The only time the result is false is if both conditions are false.

```
1 var a = true;
2 var b = false;
3 var result;
4
5 if (a) {
6     result = true;
7 }else if (b) {
8     result = true;
9 }else {
10     result = false;
11 }
12
13 if (a || b) {
14     result = true;
15 }else{
16     result = false;
17 }
```

These kind of examples also help to explain what is called short-circuit behavior in boolean operators. Take a look back at the 'and' example. If the value of `a` was false, then the first if-else structure would never even check the value `b` because it would have immediately gone to the else part. The same is true with the `&&` operator. If `a` is false, then the value of `b` doesn't matter and so it doesn't check it. For the `||` operator, if the first value is true, then the value of the second doesn't matter since only one of them has to be true for the result to be true.

It may seem like we don't need to worry about short circuit behavior as programmers, since why should we care if it checks a value or not. For the most part it doesn't matter for writing programs. However, if you make a function call in the condition of an if-statement, short-circuit behavior could matter if that function call does something in addition to being a condition.

8.3 while

The iteration control structure has to be tied to a conditional structure, so that it knows when to stop. A `while` loop initially behaves similarly to an if-statement. The condition is placed in the parenthesis immediately following the `while` keyword. If the condition is false, then execution skips the code in the braces. If the condition is true, it executes the code in the braces, but after it's done it re-evaluates the condition. It continues to execute the code in the braces, and evaluates the condition, until the condition is false.

```
1 var x = 0;
2
3 while(x < 5) {
4     x++;
5 }
```

In the above example, the variable `x` is called the loop control variable. Its initial value of 0 causes the condition `x < 5` to be true, and so it then executes the code in the braces. The only thing it does is to add 1 to the variable. It then checks the condition again, which is still true. It keeps repeating this process adding 1 to `x` each time, until it gets to 5. When `x` is 5, the condition becomes false since 5 is not less than 5. The loop executed the code a total of 5 times before it stopped.

A loop will continue to repeat itself for as long as the condition is true. If nothing occurs to cause the condition to be false, the loop will run for as

long as the program is capable: this is called an infinite loop.

```
1 var x = 0;
2 var y = 0;
3
4 while(x < 5) {
5     y++;
6 }
```

The above example will run forever, even though it might look very similar to the first example. The issue is that the variable being changed inside the loop is different than the variable being used in the condition. `x < 5` will always be true, and thus it will loop forever.

A while loop does not necessarily have to proceed in a linear fashion. How many times does the following loop repeat, and what is the final value of `x`?

```
1 var x = 1;
2
3 while(x < 32) {
4     x *= 2;
5 }
```

Consider the following loop as a brain teaser.

```
1 var x = 1;
2 var divisor = 1;
3
4 while(x < 5) {
5
6     if (x % divisor === 0){
7         x /= divisor;
8         divisor++;
9     }
10
11     x++;
12 }
```

Does this loop ever stop? If so, how many times does it loop? Just by looking at the terminating condition, and assuming it does stop eventually, what must be the final value of `x`?

8.4 for

The 'for' loop is a shorter way of writing a linear while loop. If I know ahead of time exactly how many times I want the loop to repeat, there is a regular pattern I can follow. First, create a variable that starts at zero. The loop condition is as long as that variable is less than the number of times I want it to loop. And after each loop I add 1 to the variable.

```
1 for(var x = 0; x < 5; x++) {  
2     console.log(x);  
3 }
```

When you write a for loop, the stuff inside the parenthesis is not just the loop condition. It also contains the initial value of the loop control variable, and the increment operation done after each time through the loop. In every other way this for loop works exactly the same as the while loop. It loops exactly 5 times.

Now, what numbers get printed to the console? The first time through the loop the value of `x` is 0, so 0 gets printed. The 1, 2, 3, and 4. After it prints the 4 the increment `x++` adds 1 to `x`, making it 5. But `5 < 5` is false, and so the loop doesn't repeat. Even though the final value is 5, it never gets printed.

Chapter 9

Functions

In a pure mathematical sense, a function is defined as taking an input and determining some output based on that input. A function allows us to group a piece of code that we expect to re-use many times.

9.1 Inputs and Outputs

A simple example might be a function that finds the absolute value of a number. That is, it gets rid of the negative sign if the number is negative, or else it leaves it alone.

```
1 // absolute value function
2 var absolute = function(x) {
3
4     if (x < 0) {
5         return -x; // -(-1) -> 1
6     }
7
8     return x; // 1 -> 1
9 };
10
11 var y = absolute(-3.5); // -> 3.5
```

Let's take a look at all the elements of how we create functions in JavaScript. The `var absolute` is a variable just like any other variable. It is how we are going to refer to this particular function later on in the program. The assignment `=` is simply stating that once I create the function, that variable will hold a reference to it.

Next, I use the `function` keyword to state that I am defining a new function. What follows is an open and closed set of parenthesis. This defines the

argument list to the function: the inputs to the function, and how I will refer to them from within the function. `function(x)` means that this function will have 1 input, and we will refer to that input as `x` within the function as if it is a variable. However, `x` is not accessible outside the function, only inside of it.

What follows the argument list is a set of open and closed curly brackets `{}`. All of the code between these brackets is what is run whenever the function is called. This is where I put any of the control structures or calls to other functions I wish to occur to complete the task this function is supposed to perform.

The `return` statement causes the function to stop executing the function code, and return immediately back to where the function was called. A single value can also be returned to that point as in this example. The absolute value function checks to see if the number is negative. If it is, it returns the negation of the number causing it to return a positive number. Otherwise it will return the original number, which I know must already be a positive number.

I can put a `return` statement in multiple places within the function code, but only one of them will be called during a particular call to the function since it will immediately terminate the function as soon as it's reached. So, if it returns the negation of a negative number inside the `if` statement, it will never reach the other `return` statement.

After the closing bracket of the function is a semi-colon, which is placed there because of the assignment operation that assigned this function into the variable `absolute`, just like every other assignment operation such as `var z = 4;`.

Once the function is defined, I can call the function by referring to it using the `absolute` variable. Calling the function is performed by placing an open and closed set of parenthesis with the values I wish the arguments to have inside the function, in exactly the same order as they appeared in the argument list when the function was defined. In this example I find the absolute value of `-3.5` by calling the function with that as the value of the `x` argument: `absolute(-3.5)`. This is such a useful function that it is a part of JavaScript already, and can be called by referring to `Math.abs(x)`.

9.2 Function Scope

I can put any code inside a function that I wish, including defining new variables. However, when I define a variable in a function, that variable is not directly accessible outside of the function. This is called function scope. This prevents causing confusion when using the same variable name in different functions.

```
1 // computes the approximate positive square root of a number
2 var sqrt = function(x) {
3     var guess = 0.5 * x;
4     var max_error = 0.001;
5
6     while(absolute(guess*guess - x) > max_error) {
7         guess = 0.5 * (guess + x/guess);
8     }
9
10    return guess;
11 };
12
13 var y = sqrt(2); // -> 1.4142156...
14 // exact square root of 2 is closer to 1.4142135...
```

This example shows one particular way of approximately computing the square root of a number using iteration. JavaScript has a square root function already as `Math.sqrt(x)`, but let's look at how this function works.

It begins just like the `absolute` function, by declaring a new variable `var sqrt =` to name the **function**, and an argument list `(...)` naming a single input `x`. However, this `x` is not the same as the previous `x`. The argument is only referenced inside of the function, and since this is a different function the argument can have the same name.

The first thing done when the function is called is to declare a variable `guess`, which will end up being the approximate value of the square root of the input. The initial value of `guess` is not super important, but I know that the square root of a number is less than the number itself, so starting with a guess of half the original number is a good start.

Another variable `max_error` is then declared to hold a value of how accurate I want the answer to be. I might be tempted to just put zero here, but if I did that then the function may not ever finish running! Basically, if I want a more accurate answer, the function will take longer to run since it

will have to iterate more times. So, I just pick a value of the error that is small compared to the numbers I am interested in.

These variables are not accessible outside of the function. They don't exist until the function is called, and once the function stops they in effect disappear. So, if I want a variable declared in a function to have a particular value, then I have to assign it within the function. If I want to know what the value was once the function returns, then I have to make sure there is another way to access it, such as using the `return` statement to return the value the variable contained.

A while-loop control structure is then used to implement the iteration. The loop should continue to iterate until the error is less-than or equal to the maximum error that I specified. To do that, the square of the `guess` is computed by multiplying it by itself, and then subtracting from what should be the true value of the square, which was the original number. If the guess was exactly equal to the square root of `x`, then this difference would be zero. However, it could be more or less than the true value since it is only approximate.

The difference `guess*guess - x` is then fed into the `absolute()` function from earlier because the difference could be negative, and I only care about how big it is in magnitude so that I can use the comparison operator to see if it's bigger than the maximum error I want.

Notice that I did some math inside of the argument list of the `absolute()` function. That is ok! Before `absolute` is called, all of the operations inside there are performed to result in a single number that is then assigned as the argument to the function, and only then is `absolute` actually called. Also, notice that `absolute` is really a variable declared outside of the `sqrt` function, but I am referring to it from inside the `sqrt` function. Function scope works like a one-way mirror: it can see all the variables declared outside, but no-one on the outside can see the variables that are inside.

Inside the while loop the `guess` is made better by using some mathematical magic based on the previous guess and the original number. Every time through the loop the `guess` gets closer and closer to the real deal, and eventually it's close enough that the loop condition causes the loop to stop.

After the loop is done I just need to return the value of `guess` back to where the function was called. At this point the variables `guess` and `max_error` can no longer be referenced from anywhere, but the returned

value is assigned into the variable `y` in this example, so that is where it is now stored.

9.3 Recursion

The code inside a function can call other functions, including itself. When a function calls itself, or results in another call to itself, this is called recursion. This is a form of iteration since the same code is executed repeatedly, and so must have some condition when it will stop, just like with a loop.

```
1 // computes the factorial of an integer
2 var factorial = function(n) {
3
4     if (n === 0) {
5         return 1;
6     }
7
8     return n * factorial(n - 1);
9 };
10
11 var a = factorial(5); // -> 120
```

This example shows a recursive function that computes the factorial of an integer: $n! = n(n-1)(n-2)\dots(3)(2)1$. The first thing it does is see if the argument to the function is 0, since factorial of 0 is 1. If it's not 0, then it computes the factorial of 1 less than the argument by calling itself on `n-1`, multiplies the result by `n` and returns that. It can call itself because the variable `factorial` is already defined, and the function will be defined before it gets called the first time.

Why does this work? The first time the function is called, the argument is equal to the number it's actually computing the factorial of. Say it's the number 5. We know that $5! = 5*4*3*2*1$, and we noticed that $4! = 4*3*2*1$. So we can re-write the original problem as $5! = 5 * 4!$. Translating that into code means that the factorial of a number can be computed by calling the same function on 1 less than that number, and multiplying the result by the number.

When the function is called the second time, the argument `n` has a value of 4. Technically speaking, the first call to the function hasn't even finished yet.

It's waiting for the second call of the function to return before it can multiply by it's version of `n` to return the total answer. This is ok because every time the function is called it creates an entirely new `n` to work with, so the computer doesn't get confused about which one is which. That is called a stack frame, and every recursive call creates a new frame on the stack of all the previous calls that are waiting on this particular call to the function to finish.

To complete this call, the same logic applies as the first time. $4! = 4 * 3!$. This keeps working until it gets down to zero. $1! = 1 * 0!$, but then what is $0!$? If we let the function keep doing the same thing it would think that $0! = 0 * (-1)!$, and then $(-1)! = (-1) * (-2)!$ and so on. For one thing that's just wrong. For another, it would never stop! That is what the if-statement is for: it's a terminating condition for when it reaches 0.

Once the terminating condition is reached, it returns a value without having to call itself anymore. All the previous calls to the function on the call stack that are waiting can then be completed in the reverse order. `factorial(0)` returns 1. Then `factorial(1)` returns 1. Then `factorial(2)` returns 2. `factorial(3)` returns 6. `factorial(4)` returns 24. And finally `factorial(5)` returns 120, which was the first thing to get called and is the final answer.

As an exercise, lets see what would happen if we called `factorial(-1)`? You can code it up and just see. But lets also think about it logically why what happens, happens. If the argument starts out as -1, the first thing it does is check to see if it's equal to 0. It's not, so it does the recursive call sending in `n-1`, which is -2. As stated before, this will cause an infinite recursion because the terminating condition `n == 0` will never be reached. However, it won't run forever. What will happen is that it continues to pile recursive calls onto the call stack, potentially thousands of times, but at some point the computer will say that the maximum size of the stack was reached and will terminate the program, called a stack overflow.

9.4 Exception Handling

The function call `factorial(-1)` never returns, but it doesn't run forever either. It simply exhausts the resources of the computer and crashes the program. For one, `factorial` isn't really even defined for -1. So it shouldn't have even attempted to compute it. This is an undefined scenario. The way

I might deal with that is by throwing an exception from THE function to tell whoever tried to compute $(-1)!$ that they're not suppose to do that.

```

1 // computes the factorial of an integer
2 var factorial = function(n) {
3
4     if (n < 0) {
5         throw "factorial undefined for " + n;
6     }
7
8     if (n === 0) {
9         return 1;
10    }
11
12    return n * factorial(n - 1);
13 };
14
15 var a = factorial(-1); // -> will cause an uncaught exception

```

The `throw` keyword works similarly to `return`: The function will stop executing, and a value will be associate with the termination. In this case I send a message about what happened. When something is thrown, however, it is an error, and causes what is called a program exception (something that isn't supposed to happen). If that exception is not handled by whatever called the function, then the whole program will terminate. This prevents the undefined behavior from happening unknown to anyone. The stack overflow was also an exception, but the exception I add should catch the problem sooner and give more information about what the issue is.

The way we can handle possible exceptions from a function is called a try-catch block. I put the code that might cause an exception inside the braces of the `try`. If no exception occurs, then everything proceeds as if there is no exception handling at all. However, if that code results in an exception, then execution in the try block terminates, and the catch block begins executing and the `err` argument contains the value that was specified at the `throw` statement.

```

1 try{
2     var a = factorial(-1); // -> will cause an exception
3 }catch(err) {
4
5     console.log(err); // prints "factorial undefined for -1" to the console
6
7 }

```


9.5 Side-effects

Since the code inside a function can access the variables declared outside of the function, it can also alter the values stored in those variables. Since those variables are not arguments to the function, nor are they assigned from the `return` from the function, we call this type of alteration a *side-effect* of calling the function.

```
1 // how many times the increment function was called
2 var i = 0;
3
4 // add one to i every time it's called
5 var increment = function() {
6     i = i + 1;
7
8     return i;
9 };
10
11 increment(); // i -> 1
12 increment(); // i -> 2
13 increment(); // i -> 3
14
15 var number_of_calls_so_far = i; // -> 3
```

In this example `increment` is a function that has no arguments, and does not return anything. However, when it is called it alters the variable `i` by adding 1 to its current value. Since `i` is declared outside of the function, it can be accessed directly even after the function returns.

This type of behavior is very useful and powerful, but it must be used carefully. Since side-effects can have consequences far away from where the function is actually called, as opposed to the return statement, and persist after the function returns, someone that didn't write the function may not understand what it's doing. Documentation should be used to describe what is happening if a function has side-effects.

9.6 Closure

A closure happens when one function is defined inside the scope of another function, but is still callable after the outer function has already returned. It is also a way of using the power of function side-effects without any of the dangers by containing the side effects to a particular function scope.

```
1 // increasing integers starting from 0
2 var makeIncrement = function() {
3     // how many times the increment function was called
4     var i = 0;
5
6     var increment = function() {
7         i = i + 1;
8
9         return i;
10    };
11
12    return increment;
13 };
14
15 var myIncrement = makeIncrement();
16
17 console.log(myIncrement()); // i -> 1
18 console.log(myIncrement()); // i -> 2
19
20 var anotherIncrement = makeIncrement();
21
22 console.log(anotherIncrement()); // i -> 1
23 console.log(anotherIncrement()); // i -> 2
24 console.log(anotherIncrement()); // i -> 3
25 console.log(anotherIncrement()); // i -> 4
26
27 console.log(myIncrement()); // i -> 3
```

I've taken the code for the `increment` function, and placed it inside another function called `makeIncrement`. When `makeIncrement` is called, it creates a new instance of `increment`, which is returned. However, since `increment` has to have access to the variable `i` to work, a closure is formed around that function scope so that it still has access to the version of `i` created at the same time the function was, even after `makeIncrement` has returned.

When `makeIncrement` returns the inner function, it is stored in a new variable. Whenever that version of `increment` is called it is working with its own version of `i`. I can call `makeIncrement` again, creating a new version of `increment` with its own version of `i`, independent of the other one. That is because each call to `makeIncrement` creates its own scope, and a new closure when it returns.

No matter which order the two versions of the function are called, it remembers how many times that particular version has already been called independent from the other one. However, I cannot directly access the variable

i because it was inside the function scope of `makeIncrement`. It is basically a hidden variable that can be used to keep track of some information between function calls, and so as the creator of `makeIncrement` I don't have to worry about anyone messing with that variable or making any assumptions about it.

Chapter 10

Data Structures

10.1 Objects

The 'Object' data structure is like a container that can hold named references to values, like variables. In JavaScript, everything that is not a primitive data-type is an object. Functions are also objects that can be passed around and assigned to variables and properties.

The most basic object is one that does not hold any other information, and is created by typing an open and closed brace right next to each other {}, with nothing inside.

```
1  /*
2   * Creating an empty object
3   */
4
5  // declare a variable name using 'var' keyword
6  var variableName;
7
8  // assign empty object "{}" to the variable
9  variableName = {};
10
11 // or do it in one line
12 var anotherVariable = {};
```

The named fields of an object are called properties. Properties can be defined for an object at the moment it is created in what is called an object literal. Properties created this way can be named with any valid, non-reserved name. The values of the properties can be any primitive type, including another object.

```

1  /*
2   * Adding properties when the object is created
3   */
4
5  var myObject = { // variable name, =, and open brace on same line.
6      // each property on its own line
7      // 'var' is NOT used to define properties
8
9      aNumber : 3.5, // property name, then colon, then value of property
10
11     // a comma is placed after each property until the last one
12
13     aString : "some text",
14
15     aBoolean: true,
16
17     anObject : { // we can put an object literal inside another object literal
18         anotherNumber : 98.6
19     }
20
21     // no comma after the last property
22
23 }; // a semi-colon after closing brace
24
25
26 // the properties of the object can also be defined on a single line
27 var anotherObject = {aNumber : 3, aString : "some text"};

```

Properties can be referenced in one of two different ways. If the property is a valid variable name, the property name can be placed after the variable name with a period in between.

`objectVariable.propertyName` → `value`

In the above code example, the value of a property `aNumber` of the `myObject` object would result in the value `3.5`.

`myObject.aNumber` → `3.5`

The property can also be referenced by passing a string containing only the property name into square brackets following the object. This might be useful if a variable contains a string that refers to an object property. The variable can be used in the brackets to get the value of the property.

`objectVariable ["propertyName"]` → `value`

As in the above example...

`myObject ["aNumber"] → 3.5`

Properties can be added after the object is created by referencing the new property name, and assigning a value to it.

```
1  /*
2   * Adding properties after the object is created
3   */
4
5  // start with an empty object
6  var thirdObject = {};
7
8  // properties are added when they are referenced on the object.
9
10 // referring to a property with the dot notation
11 thirdObject.aNumber = 3;
12
13 // referring to a property using the bracket notation
14 thirdObject["aString"] = "some text";
15
16 var stringVariable = "aBoolean";
17
18 // use a variable that has a string value to get the property
19 thirdObject[stringVariable] = true;
```

10.1.1 Code Exercises

1. Make an object and assign to a variable with two properties. One property named 'name' with a string value containing your name, and an 'age' property with a number value of your age.
2. Make an object and assign to a variable with two properties. An 'information' property that is itself an object structured as in exercise 1, and a 'id' property that has an integer number value (you choose the value of the number).

10.2 Arrays

Another important data structure is called an array. One can think of an array as a bunch of variables all crammed into a single name. Say we want to store 10 different numbers at a time. Without an array we would have to list out 10 different variable names. But with an array we can simply put them together in a single variable.

Lets start with an empty array. An array is started with brackets, and an empty array is two brackets [] with nothing between them.

```
1 // create an empty array
2 var emptyArray = [];
```

If we wish to start the array with something in it, we simply list those values between the brackets, separated by commas. In this example I simply put in the characters 'a', 'b', 'c', 'd'. Each individual "bin" of the array is called an element of the array. And each element has some value. In this example, the first element of the array has the value 'a'. The second element has the value 'b'.

```
var myArray = ['a', 'b', 'c', 'd'];
```

To access an element of an array, you need to specify which element you want, since there could be several. Elements are ordered by a number called an index. The index starts at 0, and increases consecutively by 1 until the last element. The index is placed in square brackets immediately following the name of the variable holding the array, and returns the value of that element.

```
arrayVariable [ index ] → value
```

The index can be thought of as the distance from the beginning of the array. If the index is 0, then we are talking about something zero distance away from the beginning, which must be the first element. So, the first element has index 0. In the example, the array variable is named `myArray`. The first element has a value of 'a', so this piece of code would return a value of 'a'.

```
myArray [ 0 ] → 'a'
```

The second element is a distance of 1 away from the beginning, so its index is 1.

```
myArray [ 1 ] → 'b'
```

An array also has a property called `length` that you can access to know how many elements are in the array. Since the index starts at 0, the index of the last element is one less than the length. The array in `myArray` has four elements, so the length of that array is 4. But the index of the last element is 3 because it is distance 3 away from the beginning of the array. The brackets have to be done last in order of operations, so any arithmetic (+, -, *, /) inside the brackets will be done first, before the index is used to look up the element.

```
myArray.length → 4
myArray.length - 1 → 3
myArray [ 3 ] → 'd'
myArray [ myArray.length - 1 ] → 'd'
```

Here is an executable example. Try writing and executing some simple code like this to get used to making and accessing an array. What value do you get when you access index of 4, which is one past the end of the array?

```
1 // array literal
2 var myArray = ['a', 'b', 'c', 'd'];
3
4 var firstElement = myArray[0];
5 // 'a'
6
7 var secondElement = myArray[1];
8 // 'b'
9
10 var nextToLastElement = myArray[myArray.length - 2];
11 // 'c'
12
13 var lastElement = myArray[myArray.length - 1];
14 // 'd'
```

10.3 JSON

JSON is a sub-set of the JavaScript language, and stands for JavaScript Object Notation. JSON is not a programming language: it cannot be used to compute anything. It only includes the idea of object, array, number, and string literals. The purpose of JSON is that perhaps you want to store or transmit text-based information around that can easily be incorporated into your program.

Since JSON has the same structure as the data in a JavaScript object, importing it is really simple and can be automated. When a JSON file is read, it can be interpreted directly into objects, arrays, numbers and strings which you can reference just like everything else in the language.

Chapter 11

Constructors

A constructor is a function that helps make a new object.

11.1 new Constructor

You will inevitably encounter constructors that require the use of the **new** keyword. The convention is that the constructor will start with a capital letter if **new** is required. As an example of how to use **new**, an empty object would be created as follows.

```
1 var emptyObject = new Object();
2 // equivalent to emptyObject = {};
```

Using that type of constructor is fairly strait-forward. You just have to remember to put a **new** in front of it. However, *writing* constructors that use **new** is outside the scope of this text. You will not need **new** to use the functional constructors we will write, and that is why I have named these constructors beginning with lower-case **make**.

11.2 Inside a Functional Constructor

A functional constructor could be as simple as returning a new object defined by an object literal.

```
1 var makePi = function () {
2
3   return {
4     value : 3.14159
5   };
6 };
```

To use it we simply call the function just like any other function. But it returns an object that has its own properties.

```
1 var myPi = makePi();
2
3 var pi = myPi.value;
4 // 3.14159
```

The constructor can add anything to the new object to give the object the desired functionality. I will use the convention of naming the object that will be returned from the constructor as `out` to indicate that anything added to that object will be exposed in the constructed object. Any other local variables will not be exposed, but will still be available to any methods which used them because it forms what is called a closure.

```
1 var makeCircle = function(radius) {
2
3     var out = {};
4
5     var area = Math.pow(radius, 2) * Math.PI;
6
7     out.radius = function() {
8         return radius;
9     };
10
11    out.area = function() {
12        return area;
13    };
14
15    return out;
16 };
```

When calling the constructor we pass in a number for what we want the radius to be, and it returns the circle object with a couple functions attached to it.

```
1 // make circle of radius = 2
2 var myCircle = makeCircle(2);
3
4 var area = myCircle.area();
5 // 12.566...
```

11.3 The specs Object

Another way of defining the argument to a constructor is to pass in a single object where the properties of that object refer to the specifications of the

object the constructor is meant to create. This argument can be named with any valid variable name as long as it is used consistently. I have chosen to name it `specs`.

```
1 var makeCircle = function(specs) {
2
3     var out = {};
4
5     var area = Math.pow(specs.radius, 2) * Math.PI;
6
7     out.radius = function() {
8         return specs.radius;
9     };
10
11    out.area = function() {
12        return area;
13    };
14
15    return out;
16 };
```

The `specs` object can be made using an object literal, and hopefully you can see immediately from reading the code what the inputs are and what they mean.

```
1 // make circle of radius = 2 using object literal for specs
2 var myCircle = makeCircle({radius: 2});
3
4 var area = myCircle.area();
5 // 12.566...
```

I have grown to like this method of passing parameters into a constructor for a couple different reasons. For one, it's easy to read what the inputs mean when the constructors appear in the code. It also makes it easy to change the meaning of the inputs without breaking code that's already been written.

As a simple example, suppose we decide to define a circle using the diameter instead of the radius, but there is already code written that defines it using the radius. We don't want to have to go back and re-write that code. If we used a `specs` object, we can just check to see which one is being used.

```
1 var makeCircle = function(specs) {
2
3     var out = {};
4
5     if ("diameter" in specs) {
```

```

6      // input used diameter , so define radius
7      specs.radius = specs.diameter / 2;
8  }
9
10     var area = Math.pow(specs.radius, 2) * Math.PI;
11
12     out.radius = function() {
13         return specs.radius;
14     };
15
16     out.area = function() {
17         return area;
18     };
19
20     return out;
21 };

```

Either radius or diameter can be used now with the same constructor.

```

1 var myCircleA = makeCircle({radius: 2});
2 var myCircleB = makeCircle({diameter: 4});
3
4 var areaA = myCircleA.area();
5 var areaB = myCircleB.area();
6 // both 12.5663...

```

The specs object is not free of issues though. For one, it may not be clear what needs to go into it when looking at the constructor itself. Incidentally it may not be clear what is in the returned object, either. One could look through the code and figure it out, but even in this simple example it may not be clear if the radius is required, or diameter, or both. This is where documentation is important. Here I just include a comment about what properties the input and output objects should have.

```

1 /**
2  * @param specs {radius:number} or {diameter:number}
3  * @returns {radius:function, area:function}
4  */
5 var makeCircle = function(specs) {
6     ...
7 };

```

Another possible issue is that any variable passed into a function retains a reference to that same object. If a variable is used to pass in the object, instead of an object literal, unpredictable things could happen if that variable is used to do other things.

```

1 var specs = {radius: 2};
2
3 var myCircleA = makeCircle(specs);
4
5 specs.radius = 3;
6
7 var myCircleB = makeCircle(specs);
8
9 var areaA = myCircleA.area();
10 // 12.5663...
11 var radiusA = myCircleA.radius();
12 // 3, ut-oh! should be 2!
13
14 var areaB = myCircleB.area();
15 // 28.274...
16 var radiusB = myCircleB.radius();
17 // 3

```

Can you see what happened? Since the `specs` object is used by the `radius` function, it will return the *current* value of that property when that function is called, which is after `specs` got changed, and not the value it had when the constructor was called originally. If we want to prevent this kind of mistake from being possible, we need to store those values in local variables which will retain their value no matter what happens to the `specs` object later on.

```

1 /**
2  * @param specs {radius:number} or {diameter:number}
3  * @returns {radius:function, area:function}
4  */
5 var makeCircle = function(specs) {
6
7     var out = {};
8     var radius;
9     var area;
10
11     if ("diameter" in specs) {
12         // input used diameter , so define radius
13         radius = specs.diameter / 2;
14     }else{
15         radius = specs.radius;
16     }
17
18     area = Math.pow(radius, 2) * Math.PI;
19
20     out.radius = function() {
21         return radius;
22     };
23

```

```
24     out.area = function() {
25         return area;
26     };
27
28     return out;
29 };
```

Now that piece of code will work as intended.

11.4 Extending Objects

It is convenient to re-used code to make an object that is related to an object that can already be made. This can be done by creating an object with one constructor, and simply adding or modifying its properties in a different constructor. Say we already have some code that constructs a Person object.

```
1 var makePerson = function(specs) {
2
3     var out = {};
4     var first = specs.firstName;
5     var last = specs.lastName;
6
7     out.getSignature = function () {
8         return first + " " + last;
9     };
10
11     return out;
12 };
```

And we've been creating Person's like this.

```
1 // use an object literal for specs input
2 var myPerson = makePerson({firstName: "John", lastName: "Doe"});
3
4 var sig = myPerson.getSignature();
5 // 'John Doe'
```

We can now write another constructor that makes use of the Person object that can already be made by simply calling that constructor, then changing the object. In this example, the `getSignature` function is redefined to include a title. Remember that functions are themselves just objects, so we can save the original function in a local variable, and then re-use it to make writing the new function a little easier.

```

1 var makeProfessional = function(specs) {
2
3     var out = makePerson(specs);
4
5     var title = specs.title;
6
7     // store the original method privately before it's overwritten
8     var super_getSignature = out.getSignature;
9
10    // overwrite method
11    out.getSignature = function () {
12        // can still re-use previous definition because we saved it
13        return super_getSignature() + ", " + title;
14    };
15
16    // add new method
17    out.earnMoney = function () {
18        return "cha-ching!";
19    };
20
21    return out;
22 };

```

A specs object can be passed in that will work with both constructors which makes life nice, and the code more consistent.

```

1 // longer object literals can be listed on new lines
2 var myPro = makeProfessional({
3     firstName: "John",
4     lastName: "Doe",
5     title: "CEO"
6 });
7
8 var sig = myPro.getSignature();
9 // 'John Doe, CEO'

```

Perhaps instead of adding new functionality to an object, we want a constructor that creates a more restricted type of object. All we have to do is create a custom `specs` object for the other constructor.

```

1 var makeDoe = function(specs) {
2
3     // add the last name
4     return makePerson({
5         firstName: specs.firstName,
6         lastName: "Doe"
7     });
8
9 };

```

Only the first name is needed because the constructor will add the last name when it calls the Person constructor.

```
1 // use an object literal for specs input
2 var myPerson = makeDoe({firstName: "Jane"});
3
4 var sig = myPerson.getSignature();
5 // 'Jane Doe'
```

11.5 Copy Constructor

Suppose we have an object already defined, but we want to make another copy of it without having to figure out what the specs object would have to be. We could add a case to handle that within the constructor itself.

11.6 Getters and Setters

11.7

11.8 Sharing Hidden Information

Hiding information is accomplished by storing values in the local variables of the constructor. These variables are not directly accessible outside of the function, unless a method is added to the `out` object to do something we want. However, there could be a time that we would like to extend an object and have access to some information defined in another constructor, but still not expose it in the resulting object.

Doing this requires a slightly different structure. Instead of having a single layer functional constructor, we need two layers. The first layer creates the object and returns the exposed object, but also an object that contains information that is meant to be hidden. I will name this function starting with the word `extend`.

The second layer hides the information that isn't supposed to be visible and looks like a normal constructor from the outside world. Even though these have to be two separate functions, we could still group them together into a single object, and have the properties of that object be the two functions. Here is an example of a Linear Congruential Generation (pseudo-random numbers).


```

1 // random integers between 0 and 100
2 var BadRandom = {
3   make : function (specs) {
4     // only return 'out' object
5     return BadRandom.extend(specs);
6   },
7
8   extend : function(specs, hide) {
9
10    var out = {}, hide = {};
11
12    hide.a = 12345;
13    hide.c = 67891;
14    hide.m = 100;
15    hide.last = specs.seed;
16
17    out.next = function () {
18      hide.last = (hide.a * hide.last + hide.c) % hide.m;
19
20      return hide.last;
21    };
22
23    // return both objects because someone
24    // might want to use or alter the hidden information
25    return {out:out, hide:hide};
26  }
27 };

```

We can use the make function property to make one of these with a seed, and see what numbers come out.

```

1 var curNum = 86;
2 var myRand = BadRandom.make({seed: curNum});
3
4 curNum = myRand.next();
5 // 61
6 curNum = myRand.next();
7 // 36
8 curNum = myRand.next();
9 // 11
10 curNum = myRand.next();
11 // 86
12 curNum = myRand.next();
13 // 61, it's repeating already, not very random

```

Testing it out we see that this is not a very good generator, but maybe we have an idea how to change some hidden variables to make it better. We can use the extend method to make a new generator but still re-use the old object.

```
1 // random integers between 0 and 100
2 var BetterRandom = {
3   make : function (specs) {
4     return BetterRandom.extend(specs).out;
5   },
6
7   extend : function(specs) {
8
9     // extend from the generator we already have
10    var random = BadRandom.extend(specs);
11
12    // make some changes to the constants
13
14    // a-1 divisible by all factors of m = 100 (2x2x5x5)
15    random.hide.a = 341; // a-1 = 2x2x5x17 -> a = 341:
16
17    // c and m relatively prime
18    random.hide.c = 143; // c = 11x13 = 143
19
20    return random;
21  }
22 };
```

Test it out to see how long it takes to start repeating the same sequence of numbers.

This example may seem silly because we could have just changed the original code. The key is that anything added to hidden can be accessed or changed by another extend function.

Chapter 12

Searching and Sorting

Index

abstraction, 19
and, 34
arguments of a function, 38
Array
 index, 50
binary tree, 3
closure, 45, 54
constructor, 53
exceptions, 43
for loop, 37
function, 38
function scope, 40
if, 29
instruction set, 18
JSON, 52
Morse code, 5
new, 53
Object
 empty object, 48
 property, 48
or, 34
recursion, 42
return statement, 39
sequence, 16
side effects of a function, 45
throw statement, 43
try-catch, 44
variable, 18
while loop, 35